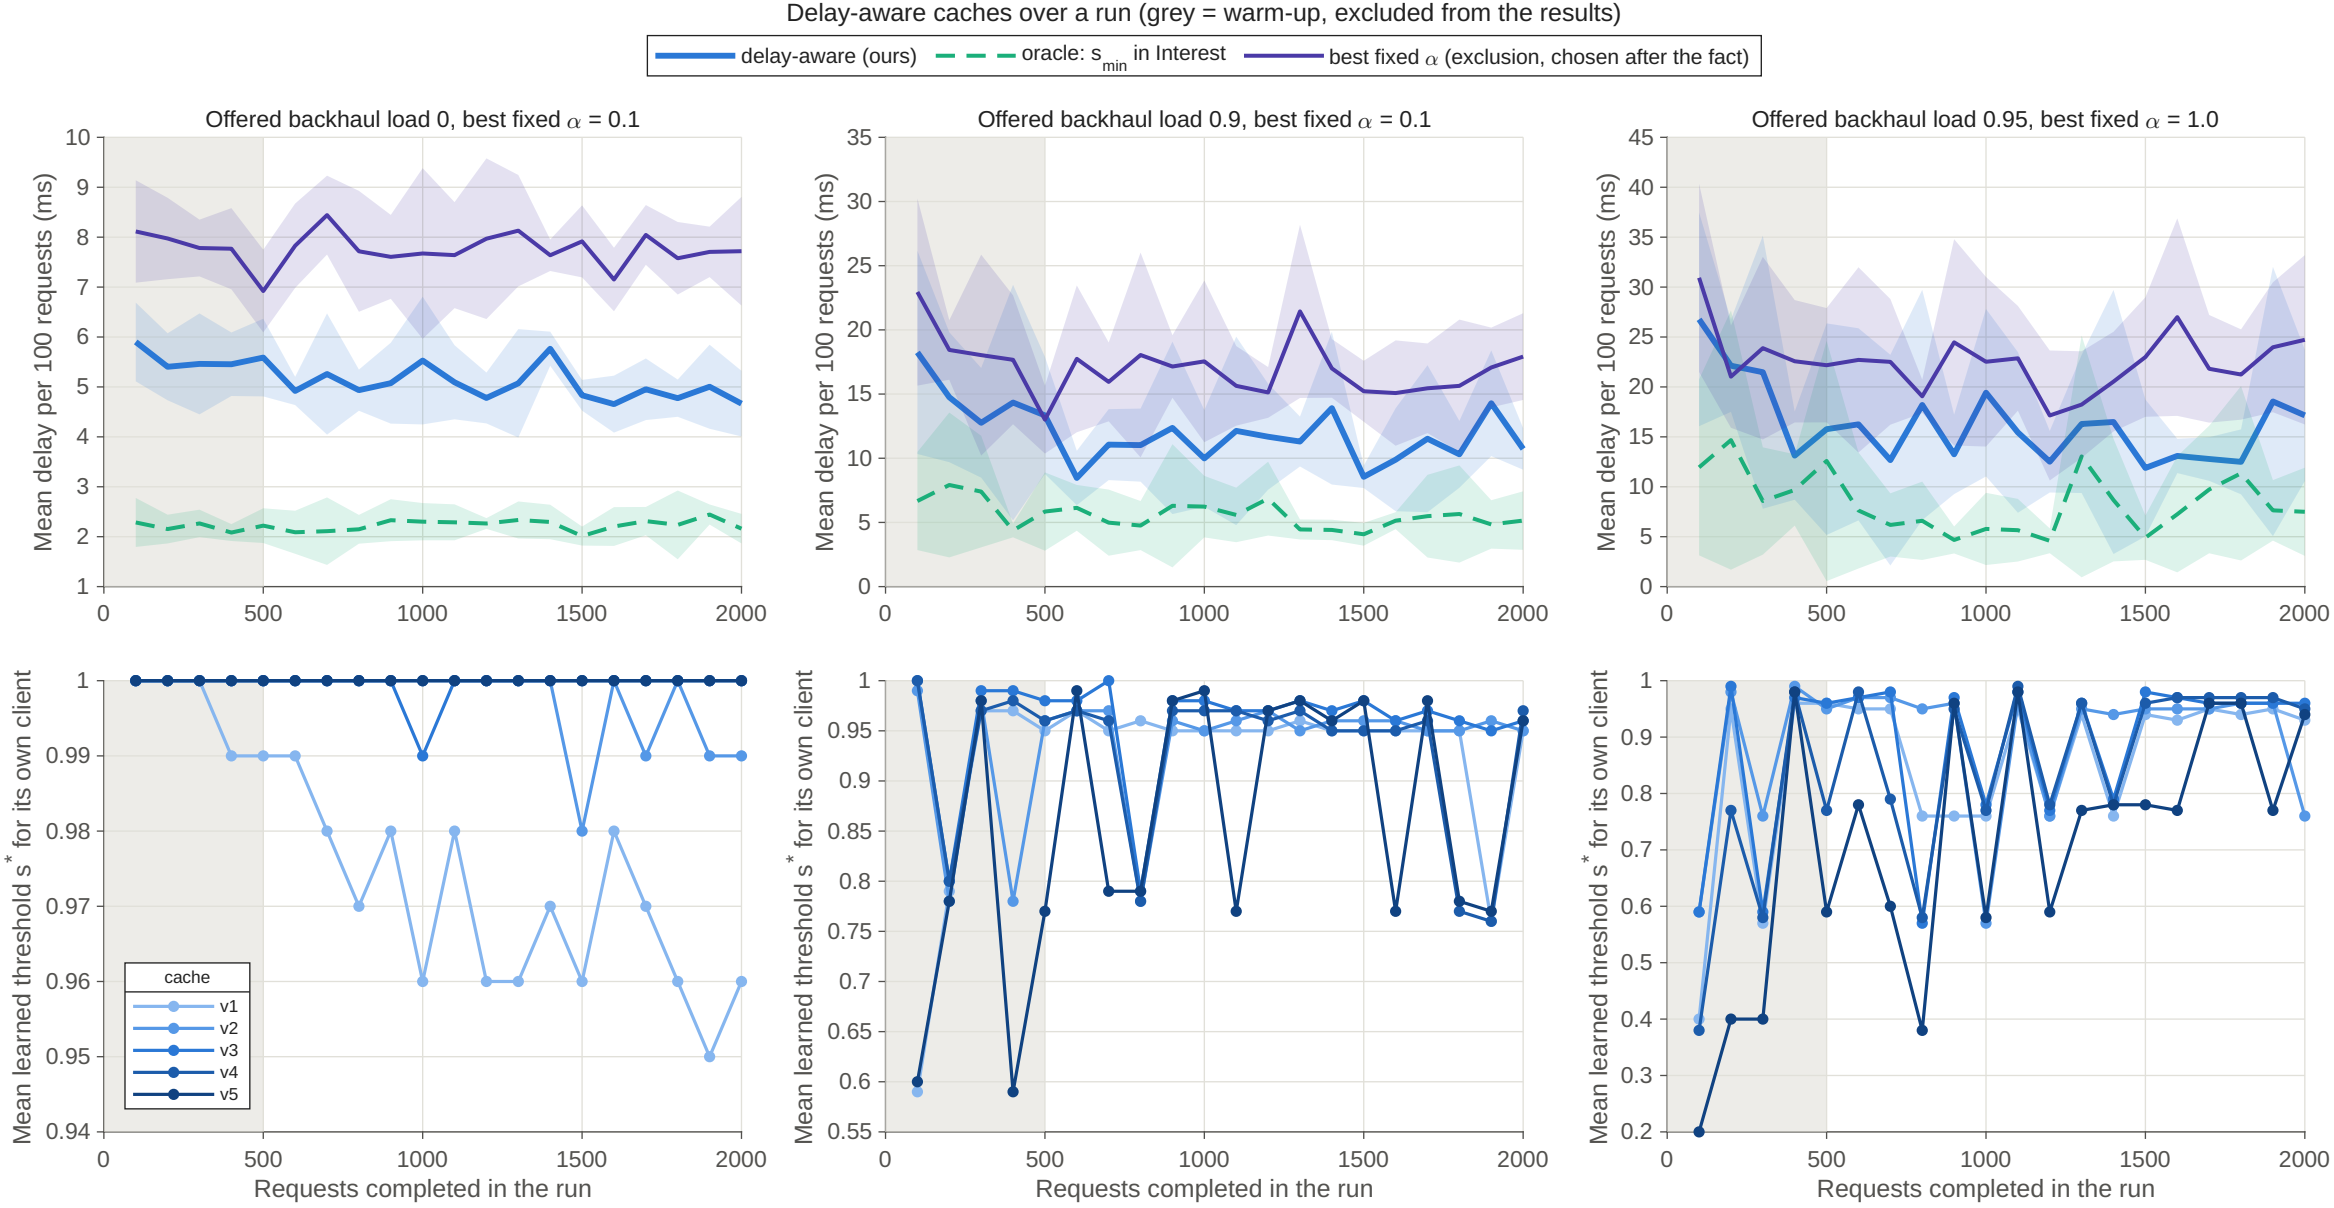
Delay-aware caches over a run (grey = warm-up, excluded from the results)
delay-aware (ours)
oracle: s_min in Interest
best fixed α (exclusion, chosen after the fact)
Offered backhaul load 0, best fixed α = 0.1
Offered backhaul load 0.9, best fixed α = 0.1
Offered backhaul load 0.95, best fixed α = 1.0
Mean delay per 100 requests (ms)
Mean learned threshold s* for its own client
Requests completed in the run
cache
v1
v2
v3
v4
v5