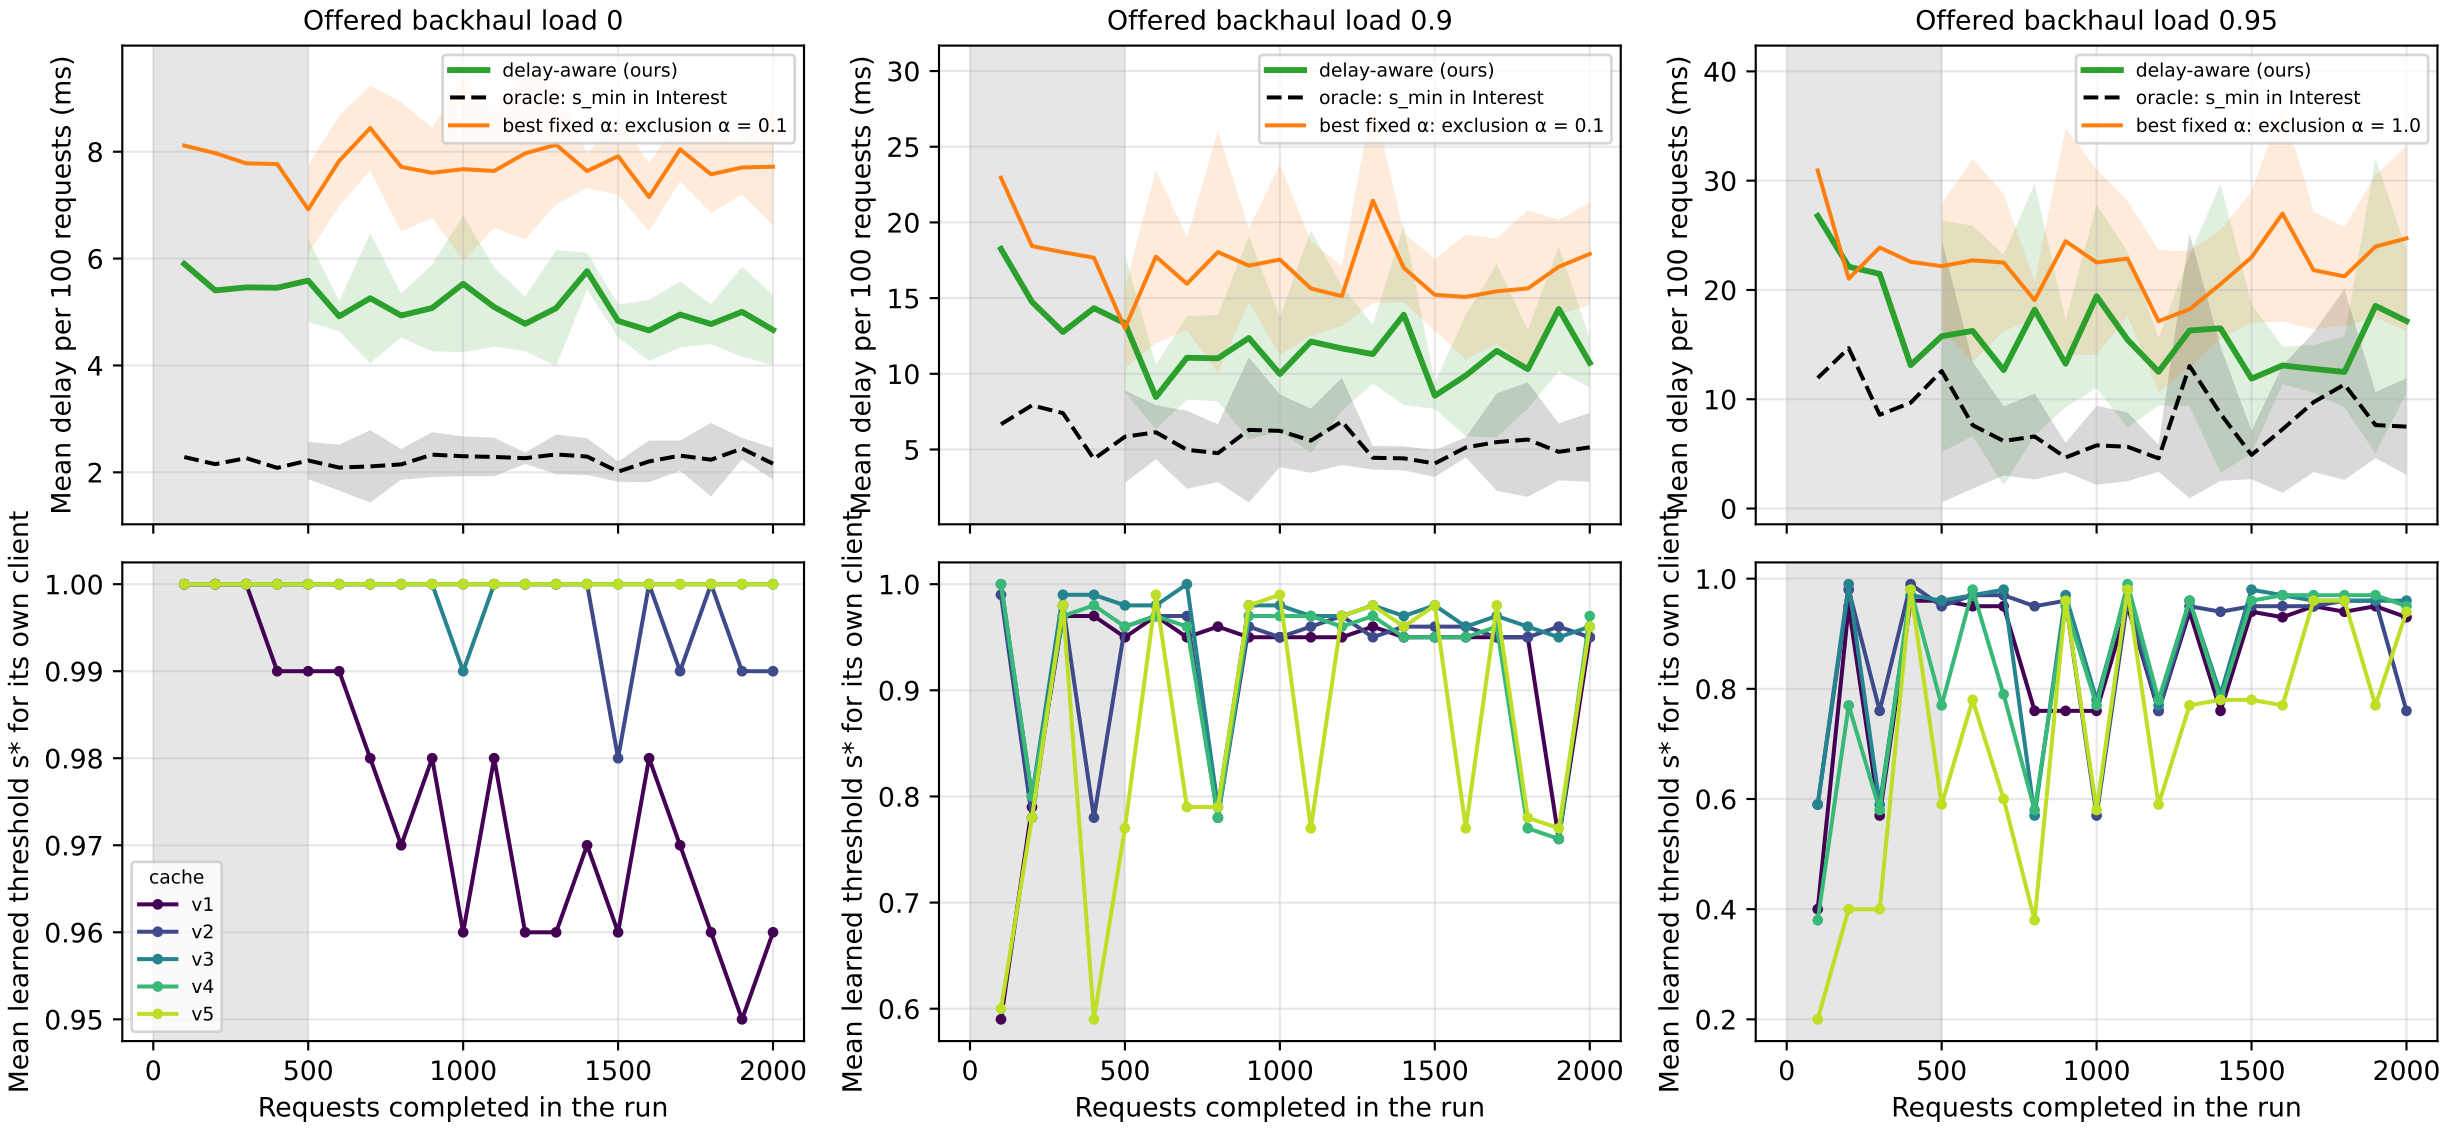
Delay-aware caches over a run (grey = warm-up, excluded from the results)
Offered backhaul load 0
Offered backhaul load 0.9
Offered backhaul load 0.95
Mean delay per 100 requests (ms)
delay-aware (ours)
oracle: s_min in Interest
best fixed α: exclusion α = 0.1
best fixed α: exclusion α = 1.0
Mean learned threshold s* for its own client
cache
v1
v2
v3
v4
v5
Requests completed in the run